

Lecture 6: Numerical Stability and R Errors

STAT 385 - James Balamuta

June 21, 2016

On the Agenda

1. Setting up Groups
 - ▶ Stable Marriage Problem
2. Numerical Stability within R
 - ▶ Observing the Variance Estimator
 - ▶ When $1 + 1 \neq 2$
3. Errors
 - ▶ Common R Errors

Grouping Algorithm

- ▶ How to Assign Groups?
 - ▶ Instructor picks randomly. Not everyone is happy.
 - ▶ Students pick themselves. Someone feels left out.
 - ▶ What about an Algorithm?

Stable Marriage Problem

- ▶ There is a fantastic TV documentary exploring algorithms on Netflix called: The Secret Rules of Modern Living: Algorithms
- ▶ Within it, the details of the Stable Marriage algorithm used in Medical School residency pairings and on Dating Websites is described.
- ▶ Video Explanation Link (Queued)

Grouping Today

- ▶ Our algorithm will try to optimize the following criteria:
 - ▶ Leadership style
 - ▶ Schedule
 - ▶ Skills
 - ▶ Project Interest

Moving along. . .

Up next, we're going to learn about implementing a **variance estimator**!

Computational Statistics and the Variance Estimator

- ▶ **Computational Statistics**, the red-headed step child between statistics and computer science, has worked time and time again to obtain an algorithm for calculating *variance*.
- ▶ Yes, **variance**.

Why is the algorithm for variance complicated?

Consider the definitions of **Mean** and the **Variance**:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$$

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

Note that the algorithm for the variance relies upon a version of the “Sum of Squares”, e.g.

$$S_{xx} = \sum_{i=1}^n (x_i - \bar{x})^2$$

That is:

$$\sigma^2 = \frac{1}{n} S_{xx}$$

Sum of Squares

Sum of Squares provides a measurement of the total variability of a data set by squaring each point and then summing them.

$$\sum_{i=1}^n x_i^2$$

More often, we use the **Corrected Sum of Squares**, which compares each data point to the mean of the data set to obtain a deviation and then square it.

$$\sum_{i=1}^n (x_i - \bar{x})^2$$

Why do we use the corrected Sum of Squares?

- ▶ When we talk about Sum of Squares it will always be the *corrected* form.
- ▶ The question for today is: **Why?**

Arithmetic Overflow

Using the initial (uncorrected) Sum of Squares definition is sure to cause an **arithmetic overflow** when working with large numbers, for example:

```
(x = (1.0024e6)^2) # Uncorrected
```

```
## [1] 1.004806e+12
```

```
(y = (1.0024e6 - 1.0000156e6)^2) # Corrected
```

```
## [1] 5685363
```

If we were to add to x, we would hit R's 32-bit integer limit (see ?integer):

```
.Machine$integer.max # Maximum integer in memory
```

```
## [1] 2147483647
```

Arithmetic Overflow - Behind the Scenes

$R > 3.0$, will try to address this behind the scenes by automatically converting the integer to a numeric with precision:

```
.Machine$double.xmax  # Maximum numeric in memory
```

```
## [1] 1.797693e+308
```

Arithmetic Overflows and Big Data

- ▶ Within *Big Data* this problem may be more transparent as the information summarized is larger.
- ▶ Thus, you may need to use an external package for very big numbers. I would recommend the following:
 - ▶ Rmpfr
 - ▶ bit64

Forms of the Variance Estimator

- ▶ Two-Pass Algorithm Form:

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

- ▶ Naive Algorithm Form:

$$\sigma^2 = \frac{\sum_{i=1}^n x_i^2 - \left(\sum_{i=1}^n x_i \right)^2 / n}{n}$$

Sum of Squares Manipulation for Naive version

I'm opting to simply show the S_{xx} modification instead of working with σ^2 since it just scales the term by $\frac{1}{n}$.

$$S_{xx} = \sum_{i=1}^n (x_i - \bar{x})^2$$

Definition

$$= \sum_{i=1}^n (x_i^2 - 2x_i\bar{x} + \bar{x}^2)$$

Expand the square

$$= \sum_{i=1}^n x_i^2 - 2\bar{x} \sum_{i=1}^n x_i + \bar{x}^2 \sum_{i=1}^n 1$$

Split Summation

$$= \sum_{i=1}^n x_i^2 - 2\bar{x} \sum_{i=1}^n x_i + \underbrace{n\bar{x}^2}_{\sum_{i=1}^n c = n \cdot c}$$

Separate the summation

$$= \sum_{i=1}^n x_i^2 - 2\bar{x} \left[n \cdot \frac{1}{n} \right] \sum_{i=1}^n x_i + n\bar{x}^2$$

Multiple by 1

Sum of Squares Manipulation for Naive version - Cont.

$$S_{xx} = \sum_{i=1}^n x_i^2 - 2\bar{x} \left[n \cdot \frac{1}{n} \right] \sum_{i=1}^n x_i + n\bar{x}^2$$

Multiple by 1

$$= \sum_{i=1}^n x_i^2 - 2\bar{x} n \cdot \underbrace{\left[\frac{1}{n} \sum_{i=1}^n x_i \right]}_{=\bar{x}} + n\bar{x}^2$$

Group terms for mean

$$= \sum_{i=1}^n x_i^2 - 2\bar{x} n\bar{x} + n\bar{x}^2$$

Substitute the mean

$$= \sum_{i=1}^n x_i^2 - 2n\bar{x}^2 + n\bar{x}^2$$

Rearrange terms

$$= \sum_{i=1}^n x_i^2 - n\bar{x}^2$$

Simplify

Implementing Naive Variance

```
var_naive = function(x){  
  n = length(x)           # Obtain the length  
  sum_x = 0                # Storage for Sum of X  
  sum_x2 = 0               # Storage for Sum of X^2  
  for(i in seq_along(x)){ # Calculate sums  
    sum_x = sum_x + x[i]  
    sum_x2 = sum_x2 + x[i]^2  
  }  
  
  # Compute the variance  
  v = (sum_x2 - sum_x*sum_x/n)/n  
  return(v)  
}
```

Implementing Two-Pass Variance

```
var_2p = function(x){  
  n = length(x)           # Length  
  mu = 0; v = 0           # Storage for mean and var  
  
  for(i in seq_along(x)){ # Calculate the Sum for Mean  
    mu = mu + x[i]  
  }  
  
  mu = mu / n             # Calculate the Mean  
  
  for(i in seq_along(x)){ # Calculate Sum for Variance  
    v = v + (x[i] - mu)*(x[i] - mu)  
  }  
  
  v = v/n                 # Calculate Variance  
  return(v)               # Return  
}
```

Calculations

```
set.seed(1234) # Set seed for reproducibility  
x = rnorm(2e6, mean = 1e20, sd = 1e12)
```

```
(method1 = var_naive(x))
```

```
## [1] 1.318357e+27
```

```
(method2 = var_2p(x))
```

```
## [1] 1.001425e+24
```

```
(baser = var(x)*((2e6)-1)/(2e6))
```

```
## [1] 1.001425e+24
```

```
all.equal(method1, method2)
```

```
## [1] "Mean relative difference: 0.9992404"
```

R's Implementation

R opts to implement this method using a two-pass approach.

- ▶ Check out the source [here](#)
- ▶ There are quite a few papers on this topic going considerably far back. See Algorithms for Computing the Sample Variance: Analysis and Recommendations (1983)

$$1 + 1 \neq 2$$

Computers in all their infinite wisdom and ability are not perfect. One of the particularly problematic areas of computers is handling **numeric** or **float** data types. Consider:

```
x = 0.1
x = x + 0.05
x
```

```
## [1] 0.15
```


```
if(x == 0.15){
  cat("x equals 0.15")
} else {
  cat("x is not equal to 0.15")
}
```

```
## x is not equal to 0.15
```

Why isn't x equal to 0.15!?

THAT'S NOT HOW THIS WORKS

**THAT'S NOT HOW ANY OF
THIS WORKS**



Enter: Numerical Stability

In essence, R views the two numbers differently due to rounding error during the computation:

```
sprintf("%.20f", 0.15) # Formats Numeric
```


```
## [1] "0.149999999999999999445"
```

```
sprintf("%.20f", x)
```

```
## [1] "0.1500000000000000002220"
```

ϵ neighborhood

Specifically, we are hitting a machine tolerance fault given by an ϵ neighborhood.



The value of the ϵ is given by:

```
.Machine$double.eps
```

```
## [1] 2.220446e-16
```

This gives us the ability to compare up to $1e-15$ places accurately.

```
sprintf("%.15f", 1 + c(-1,1)*.Machine$double.eps)
```

```
## [1] "1.000000000000000" "1.000000000000000"
```


Discrete Solution Check

To get around rounding error between two objects, we add a tolerance parameter to check whether the value is in the ϵ neighborhood or not.

```
all.equal(x, 0.15, tolerance = 1e-3)
```

```
## [1] TRUE
```

Discrete Solution Check

Since `all.equal` may not strictly return `TRUE` or `FALSE`, it is highly advisable to wrap it in `isTRUE()`, e.g.

```
isTRUE(all.equal(x, 0.15))
```

```
## [1] TRUE
```

Thus, in an `if` statement, you would use:

```
if(isTRUE(all.equal(x, 0.15))){  
  cat("In threshold")  
} else {  
  cat("Out of threshold")  
}
```

```
## In threshold
```

Bad Loop

To magnify the issue consider a loop like so:

```
inc_value = 360L / 14L # Value to increment  
i = 0 # Increment storage  
while(i != 360){ # Loop  
    i = i + inc_value # Add values  
}
```

After 14 iterations, the loop should complete, but it does *not*! In fact, this loop will go onto infinity.

Good Loop

To fix the looping issue, we opt to always stick with *integers* as counters

```
inc_value = 360 / 14 # Value
i = 0L          # Integer
o = 0          # Numeric
while(i != 14L){
    o = o + inc_value # Sum
    i = i + 1L        # Increment loop
}
i
```

```
## [1] 14
```

Summary

- ▶ Statistics and Computer Science rely on each greatly in this Brave New World of Data Science.
- ▶ When working with big numbers, understand that arithmetic overflow is a reality.
- ▶ **Never, ever, ever** use a floating-point representation as an incrementor for a loop.
 - ▶ Always use an `integer` for an incrementor and then convert it to a `numeric` within a function.
- ▶ This topic **will** come up again when we switch to using `Rcpp`.

Coming up next. . . .

Common *R* errors. . .

Checking for Equality vs. Assignment

The most common error by far that affects programmers is making an assignment when trying to check for equality (and vice versa)

```
# Assigning in `if`  
if(x = 42) { cat("Life!") }  
## Error: unexpected '=' in "if(x ="
```

```
# Correct  
if(x == 42) { cat("Life!") }
```

```
# Equality Check instead of Assignment  
x == 42  
## No Error, but prints `TRUE` or `FALSE`
```

```
# Correct  
x = 42
```

if vectorization usage

As emphasis on vectorization grows, there is a tendency to compare two vectors using the default `if()` instead of `ifelse()`

```
x = 1:5
y = 2:6
if(x > y){ T }
## Warning messages:
## In if (x > y) { : the condition has length > 1 and only
# Correct
ifelse(x > y, T, F)
```


Vector Recycling

Sometimes the length of vectors are not equal or the data does not divide evenly or oddly when perform a vectorized computation.

```
x = 1:5
```

```
y = 2:3
```

```
x + y
```

```
## Warning message:
```

```
## In x + y : longer object length is not a multiple  
## of shorter object length
```

```
# Correct
```

```
x = 1:4
```

```
y = 2:3
```

```
x + y
```

```
# Repeats y twice
```

```
# 1 + 2, 2 + 3, 3 + 2, 4 + 3
```

Mismatched curly brackets {} or parentheses ()

Often it is ideal to use parentheses for order of operations or curly brackets {}, though this sometimes causes a mismatch.

```
2*(x + y)
```

```
## Error: unexpected ')' in "2*(x + y))"
```

```
# Corrected
```

```
2*((x + y))
```

No Multiplication

When working on computations, sometimes we just “slip” and opt not to write a multiplication sign thinking the interpreter can understand the context.

```
2x+4
```

```
## Error: unexpected symbol in "2x"
```

```
# Correct
```

```
2*x + 4
```

Manual Data Entry

Sometimes it's easier as we'll see next week to manually enter data. The issue with this is sometimes you forget simple things like a ,.

```
c(1, 2 3, 4)
## Error: unexpected numeric constant in "c(1,2 3"

# Correct
c(1, 2, 3, 4)
```

Strings in character values

At times, there may come a need to place a quotation inside of a string. To do this, requires using an escape character \ or using ' ' instead.

```
"toad"princess"  
## Error: unexpected symbol in "\"toad"princess"  
  
# Corrected  
"toad\"princess"  
'toad"princess'
```

Handling Missing Value Operations

The NA character indicates the presence of a missing value. These missing values can play havoc with computations.

```
x = c(1,NA,2)
3 + x
# No Error, but: [1] NA

sum(x)
# No Error, but: [1] NA

# Corrected
1 + na.omit(x)      # Deletes NA
sum(x, na.rm = T)  # Removes NA inside function
```

Finiteness of Values

R can have some funky finiteness problems due to how NA values are created.

```
x = c(NA, -Inf, Inf, NaN)
is.na(x)
# No error, but: [1] TRUE FALSE FALSE TRUE

is.infinite(x)
# No error, but: [1] FALSE TRUE TRUE FALSE

# Correct
is.finite(x)
# [1] FALSE FALSE FALSE FALSE
```

Summary of Errors

- ▶ There are many odd errors that *R* can cause.
- ▶ StackOverflow is a great community to ask questions about errors.
- ▶ When dealing with NA values, be on your guard!